

# Threads & Tasks: Executor Framework

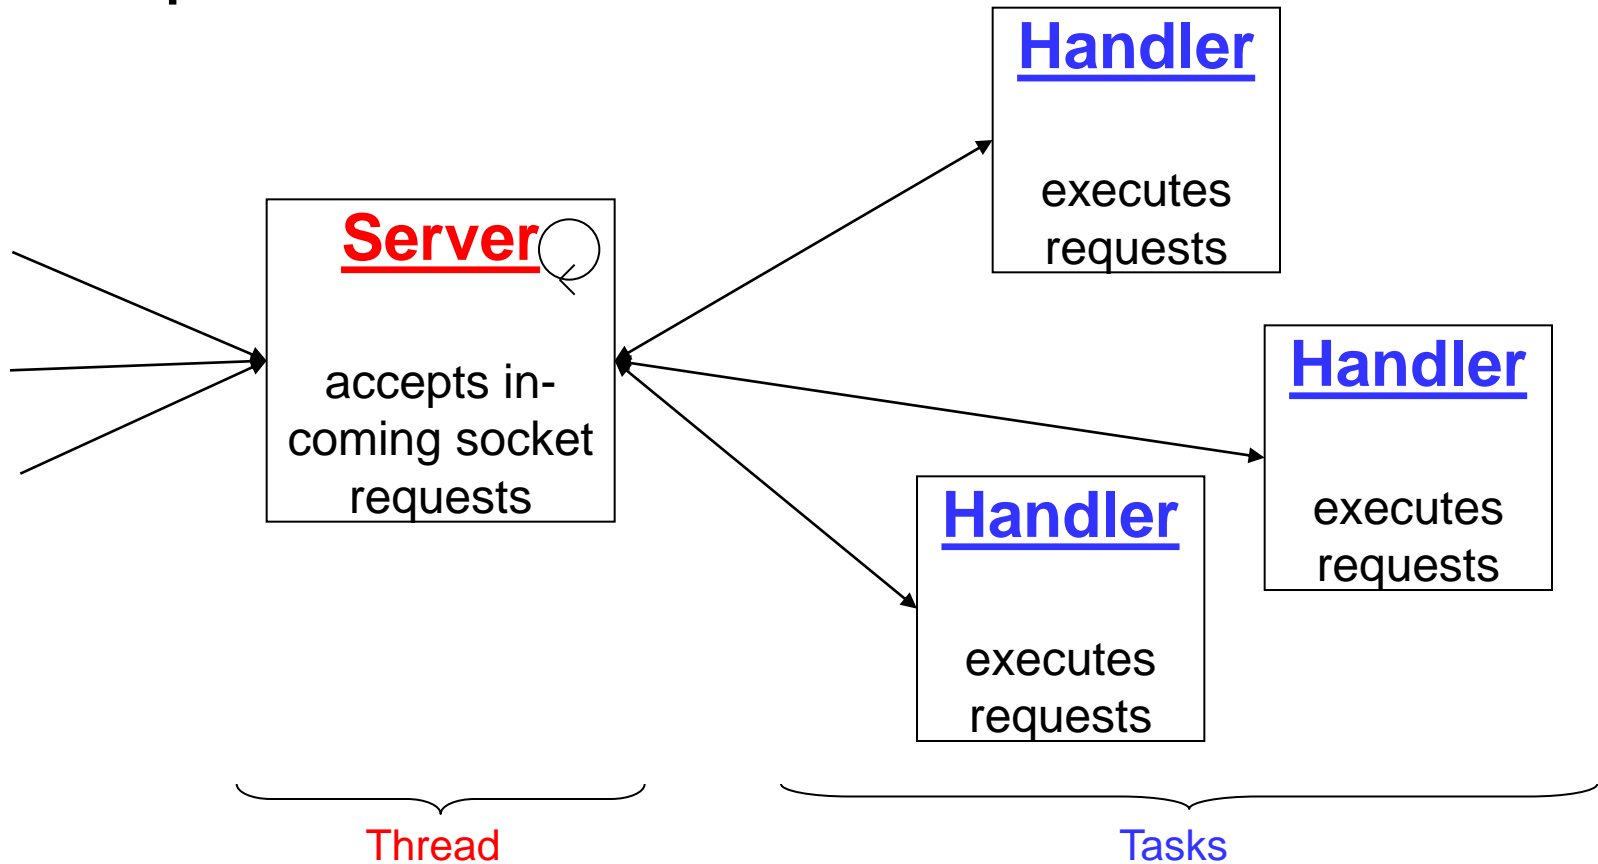
- **Introduction & Motivation**
  - WebServer
- **Executor Framework**
- **Callable and Future**

# Threads & Tasks

- **Motivations for using threads**
  - Actor-based
    - Goal: Create an autonomous, active object which can react on events
    - **Example: swing event handler thread**
  - Task-based
    - Task = Logical unit of work
    - Goal: Asynchronous execution of a task
    - **Example: mandelbrot slice task**
- **Task based concurrency**
  - Simplifies program organization
  - Facilitates error recovery by providing natural transaction boundaries
  - Promotes concurrency by providing a natural structure for parallelizing work

# Threads & Tasks

- Example: Webserver



# Introduction: Implementation of a WebServer1

```
public class webServer1 {  
    public static void main(String[] args) throws IOException {  
        ServerSocket serverSocket = new ServerSocket(80);  
        while(true) {  
            Socket s = serverSocket.accept();  
            handleRequest(s);  
        }  
    }  
}
```

- **Single-Threaded Server:**  
**Tasks are executed within the server thread**
  - Poor performance, can only handle one request at a time
  - While server is processing, new requests must wait
  - Possible if request processing is very fast

# Introduction: Implementation of a WebServer2

```
public class webServer2 {  
    public static void main(String[] args) throws IOException {  
        ServerSocket serverSocket = new ServerSocket(80);  
        while(true) {  
            final Socket s = serverSocket.accept();  
            Thread t = new Thread(){  
                public void run(){ handleRequest(s); }  
            };  
            t.start();  
        }  
    }  
}
```

- **Explicitly creating threads for tasks**  
**Each task is executed in its own thread**
  - handleRequest must be thread-safe
  - Excessive thread creation: Scheduling overhead / memory consumption

# Introduction: Implementation of a WebServer3

```
public class webServer3 {  
    public static void main(String[] args) throws IOException {  
        final ServerSocket serverSocket = new ServerSocket(80);  
        for (int i = 0; i < 10; i++) {  
            Thread t = new Thread() {  
                public void run() {  
                    while (true) {  
                        try { handleRequest(serverSocket.accept()); }  
                        catch (IOException e) { /* ... */ }  
                    }  
                }  
            };  
            t.start();  
        }  
    }  
}
```

# Introduction: Implementation of a WebServer3

- **Disadvantages**

- Maybe incorrect, `ServerSocket.accept()` is *not* documented as threadsafe
- No flexibility, i.e. exactly 10 threads
  - No creation of new threads
  - No deletion of unused threads
- All threads are pre-created, no lazy allocation
- No life-cycle management, i.e. "pool" cannot be stopped
- No error handling, if an exception is not caught, thread terminates silently
- No flexibility in the order of pending requests
  - They are stored in the Server-Socket queue
- No full control over size of queue for pending requests
  - backlog parameter only specifies *maximum* length of the queue

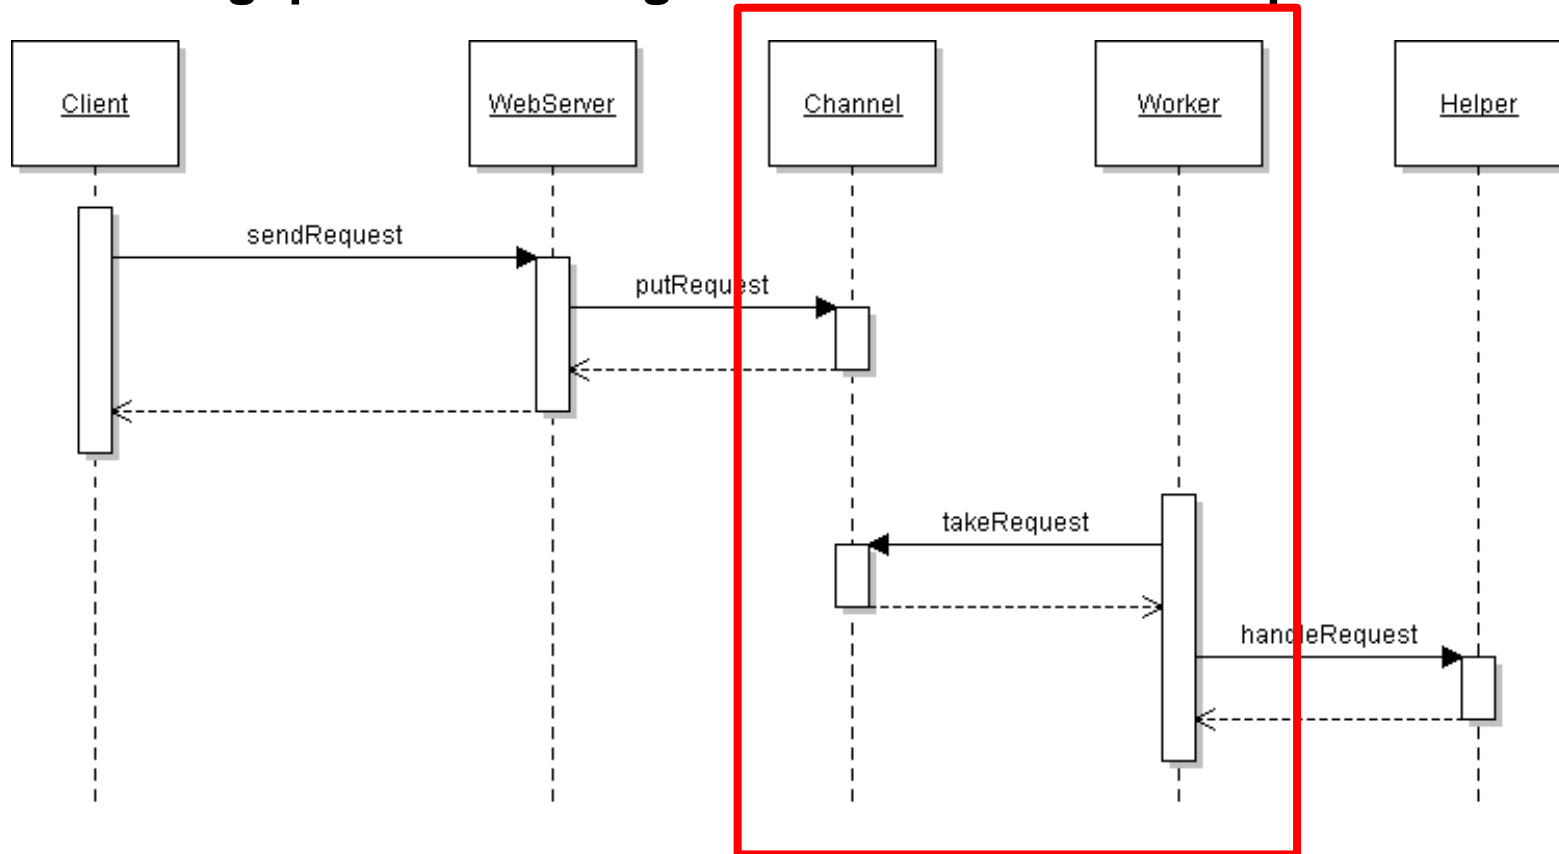
# Threads & Tasks: Executor Framework

- Introduction & Motivation
- **Executor Framework**
- Callable and Future



# Executor Framework

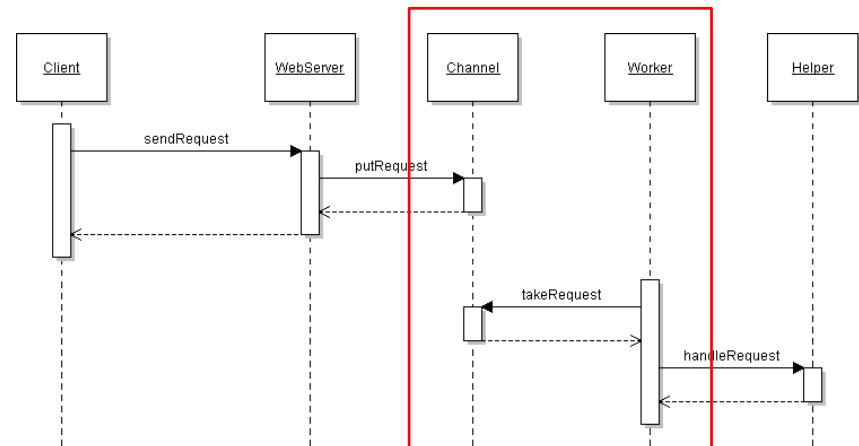
- Fills the gap between single-threaded and thread-per-task



# Executor Framework

- **Participants**

- Worker: One thread is used to execute many unrelated tasks
  - These threads are called worker threads / background threads
  - May be organized in a thread pool (if more than one thread is used)
  - May provide a flexible thread management
- Channel: A buffer which holds pending requests
  - May be bounded
  - Implements the producer-consumer pattern



# Executor Framework

- **Executor = Channel + Worker**

```
public interface Executor {  
    void execute(Runnable command);  
}
```

- Decouples task submission from task execution
- Executes the given command at some time in the future
- The command may be executed
  - in a new thread / in a pooled thread / in the calling thread

- **Runnable = Task**

```
public interface Runnable {  
    void run();  
}
```

- Limitation:
  - Method *run* cannot return a result (results are placed in shared fields)
  - Method *run* cannot declare a checked exception

# Introduction: Implementation of a WebServer4

```
public class webServer4 {  
    public static void main(String[] args) throws IOException {  
        Executor exec = new ThreadPoolExecutor(10);  
        ServerSocket serverSocket = new ServerSocket(80);  
        while(true){  
            final Socket s = serverSocket.accept();  
            Runnable task = new Runnable(){  
                public void run(){ handleRequest(s); }  
            };  
            exec.execute(task);  
        }  
    }  
}
```

## Executor: Implementation

```
class ThreadPoolExecutor implements Executor {
    private final BlockingQueue<Runnable> queue
        = new LinkedBlockingQueue<Runnable>();

    public void execute(Runnable r) { queue.offer(r); }

    public ThreadPoolExecutor(int nrThreads) {
        for (int i = 0; i < nrThreads; i++) { activate(); }
    }

    private void activate() {
        new Thread(new Runnable() {
            public void run() {
                try {
                    while (true) { queue.take().run(); }
                } catch (InterruptedException e) { /* die */ }
            }
        }).start();
    }
}
```

# Executor: Simple implementations

- **DirectExecutor: Synchronous execution (in calling thread)**

```
class DirectExecutor implements Executor {  
    public void execute(Runnable r) { r.run(); }  
}
```

- With this executor Webserver4 = Webserver1

- **Thread per task executor**

```
class ThreadPerTaskExecutor implements Executor {  
    public void execute(Runnable r) {  
        new Thread(r).start();  
    }  
}
```

- With this executor Webserver4 = Webserver2

# Executor: Advanced Implementations

- **Execution Policies**

- Execution order of submitted tasks (FIFO, LIFO, Priority Queue)
  - More control as with Java's setPriority
- Number of threads which execute concurrently
- Maximal size of queue with pending tasks
- Actions taken before / after task execution (startup/cleanup)

- **Factory: `java.util.concurrent.Executors`**

- Provides several implementations for thread pools which implement different policies
- All factory methods return an executor which implements the `ExecutorService` interface

# Executor: Advanced Implementations

- **Executors.newFixedThreadPool**
  - Threads are created up to a fixed number
  - Threads which die due to an unexpected exception are replaced
- **Executors.newCachedThreadPool**
  - Creates new threads as needed, reusing previously constructed threads if they are available
- **Executors.newSingleThreadExecutor**
  - Uses single worker thread
  - Worker thread is replaced if an unexpected exception occurs
- **Executors.newScheduledThreadPool**
  - Creates a ScheduledExecutorService which supports
    - Periodic tasks (scheduleAtFixedRate / scheduleWithFixedDelay)
    - Delayed tasks (schedule)



# ThreadFactory

- **Some factory methods on j.u.c.Executors take a ThreadFactory**

```
public interface ThreadFactory {  
    Thread newThread(Runnable r);  
}
```

- **Enables applications to use**
  - special Thread subclasses
  - priorities
  - custom named threads
  - daemon flag

# Executors: Design Considerations

- **Identity**
  - Different tasks are executed with the same thread  
=> thread-specific contextual control techniques are difficult
    - Use of ThreadLocals
    - Use of security contexts
- **Queuing**
  - Runnable tasks that are waiting in queues do not run  
=> if a currently running task blocks waiting for a condition produced by a task still waiting in a queue => system may freeze up
    - Use enough worker threads
    - Restrict dependencies
    - Create custom queues that understand the dependencies  
=> Fork-Join Framework

# ExecutorService: Executor Life-Cycle

- **ExecutorService: provides life-cycle management methods**

```
interface ExecutorService extends Executor {  
    void shutdown();  
    List<Runnable> shutdownNow();  
    boolean isShutdown();  
    boolean isTerminated();  
    boolean awaitTermination(long timeout, TimeUnit unit)  
        throws InterruptedException;  
}
```

- Supports shutdown methods
  - shutdown: Graceful shutdown: finish pending tasks, do not accept new ones
  - shutdownNow: Abrupt shutdown: running tasks are interrupted, returns list of tasks that were not started
- awaitTermination: awaits until executor service is terminated



# Executors and JMM

- **Memory consistency effects**
  - Actions in a thread prior to submitting a Runnable object to an **Executor** **happen-before** its execution begins (possibly in another thread)
  - Actions in a task which is executed by a **SingleThreadExecutor** **happen-before** actions executed in subsequent tasks (even if the subsequent task is executed by another thread due to an exception)

# Threads & Tasks: Executor Framework

- Introduction & Motivation
- Executor Framework
- **Callable and Future**

# Result-bearing tasks: Callable & Future

- **Callable: Task with a result / exception**

```
interface Callable<V> {  
    V call() throws Exception;  
}
```

- **Future: represents a future result of a task**

```
interface Future<V> {  
    boolean cancel(boolean mayInterruptIfRunning);  
    boolean isCancelled();  
    boolean isDone();  
    V get() throws InterruptedException, ExecutionException,  
        CancellationException;  
    V get(long timeout, TimeUnit unit)  
        throws InterruptedException, ExecutionException,  
        CancellationException, TimeoutException;  
}
```


# Callable & Future

- **Submitting tasks**

```
interface ExecutorService {  
    // ... Lifecycle methods  
    Future<?> submit(Runnable task);  
    <T> Future<T> submit(Callable<T> task);  
    <T> Future<T> Submit(Runnable task, T result);  
  
    <T> List<Future<T>> invokeAll(  
        Collection<? extends Callable<T>> tasks)  
        throws InterruptedException;  
  
    <T> T invokeAny(  
        Collection<? extends Callable<T>> tasks)  
        throws InterruptedException, ExecutionException;  
}
```

# Callable & Future

- **States of a Task**



- **get()**

- If completed: returns immediately  
(returns result or throws `ExecutionException`)
- If not completed: method call blocks
  - If terminates regularly => result
  - If terminates with an exception => `ExecutionException`
  - If cancelled => `CancellationException`
  - If thread calling `get` was interrupted => `InterruptedException`



# CompletionService

- **CompletionService = Executor + BlockingQueue**
  - Decouples production of new tasks from the consumption of the results of completed tasks
    - Producers submit tasks for execution
    - Consumers take completed tasks and process their results

```
interface CompletionService {  
    Future<V> poll(); // returns available future or null  
    Future<V> poll(long timeout, TimeUnit unit) throws IE;  
    Future<V> take() throws IE; // waits for a future  
    Future<V> submit(Callable<V> task);  
    Future<V> submit(Runnable task, V result);  
}
```

- take retrieves and removes the next completed task, potentially waiting
- **ExecutorCompletionService**
  - Uses a separate Executor to schedule the tasks

# CompletionService: Sample

- **Solver**

- Method solve takes a list of solvers (of type Callable<Result>)
- Result of first solver is returned

```
void solve(Executor e, Collection<Callable<Result>> solvers)
    throws InterruptedException {
    CompletionService<Result> cs = new ExecutorCompletionService<>(e);
    int n = solvers.size();
    List<Future<Result>> futures = new ArrayList<Future<Result>>(n);
    Result result = null;
    try {
        for (Callable<Result> s : solvers) futures.add(cs.submit(s));
        ...
    }
```

## CompletionService: Sample

```
for (int i = 0; i < n; ++i) {  
    try {  
        Result r = cs.take().get();  
        if (r != null) { // solver may return null  
            result = r;  
            break;  
        }  
    } catch (InterruptedException ignore) {}  
}  
}  
finally {  
    // cancel all pending solvers  
    for (Future<Result> f : futures) f.cancel(true);  
}  
  
if (result != null) use(result);  
}
```

# Callable & Future and JMM

- **Memory consistency effects**
  - Actions in a thread prior to the submission of a Runnable or Callable task to an ExecutorService **happen-before** any actions taken by that task
  - Any actions taken by a Runnable or Callable task executed by an ExecutorService **happen-before** the result is retrieved via Future.get()